Resident-expert lookahead: Code generation

Ryzen 5 7600X / Radeon RX 7900 XTX / IQ1_M / HIP

4 of 12 A/B pairs match output, MTP work and prior requests | exploratory

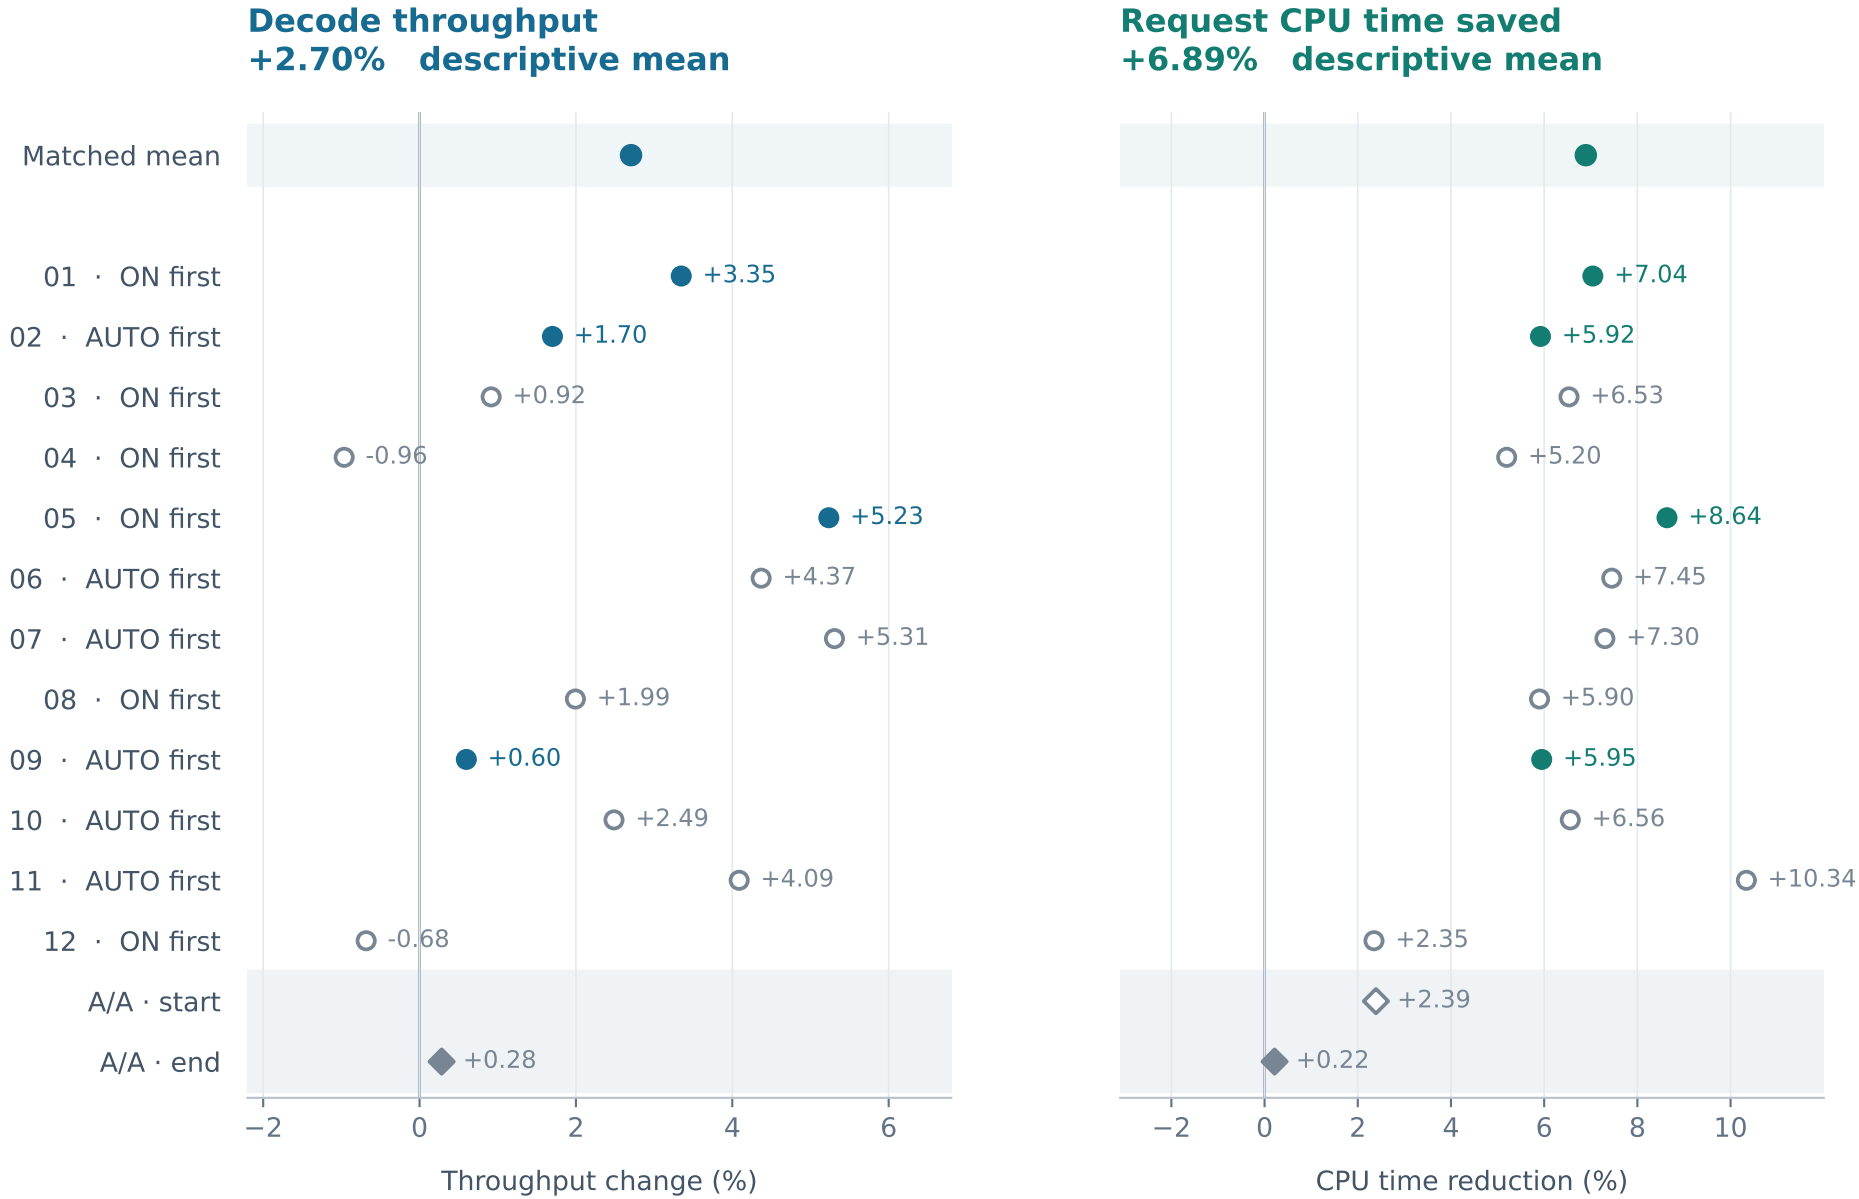

Matching work

Different work; excluded from matched estimate

Both engines ON

600-token warmup + 600-token output per workload · 128K allocated capacity · same executable

A/B: automatic vs forced ON. A/A: second ON vs first ON. English interval: 95% paired bootstrap.